

MelodyUD: Platform Design Inspired by Spotify's Business Model

Exploring Database Structures for Scalable Music Platforms

Authors:

Brayan Yate & Holman Alvarado

Course: Databases II

Year: 2025

Business Model

Spotify's Freemium Model

- **Free Users:** Ad-supported, limited skips
- **Premium Users:** Ad-free, Hi-Fi audio, offline listening
- **Revenue Streams:**
 - Subscriptions (main)
 - Advertisements (supporting growth)

Functional Requirements

- **Account Management:** Email/social sign-in, profile updates
- **Content Ingestion:** Upload tracks, enrich with ML
- **Playback & Delivery:** Adaptive bitrate, caching
- **Discovery:** Search + Recommendations
- **Advertising:** Campaigns, real-time targeting
- **Analytics:** Real-time dashboards for creators
- **Governance:** Tokenization, audit logs

Non-Functional Requirements



- **Performance:**
 - Playback start $\leq 300\text{ms}$
 - Search $\leq 150\text{ms}$
- **Scalability:**
 - $\geq 1\text{B}$ users, 20M streams
 - 10TB daily audio ingest
- **Availability:**
 - $\geq 99.95\%$ uptime
- **Maintainability:**
 - OpenAPI docs, microservices $< 3\text{K LOC}$

User Roles



- **Listener:** Play, search, create playlists
- **Creator:** Upload, view analytics, promote
- **Advertiser:** Create and monitor campaigns
- **Engineer:** Ensure uptime, scale, trace errors
- **Compliance Officer:** Enforce GDPR, log retention

Some User Stories

Listener: Create Account

-  Sign up via email or social login
-  Credentials hashed + MFA

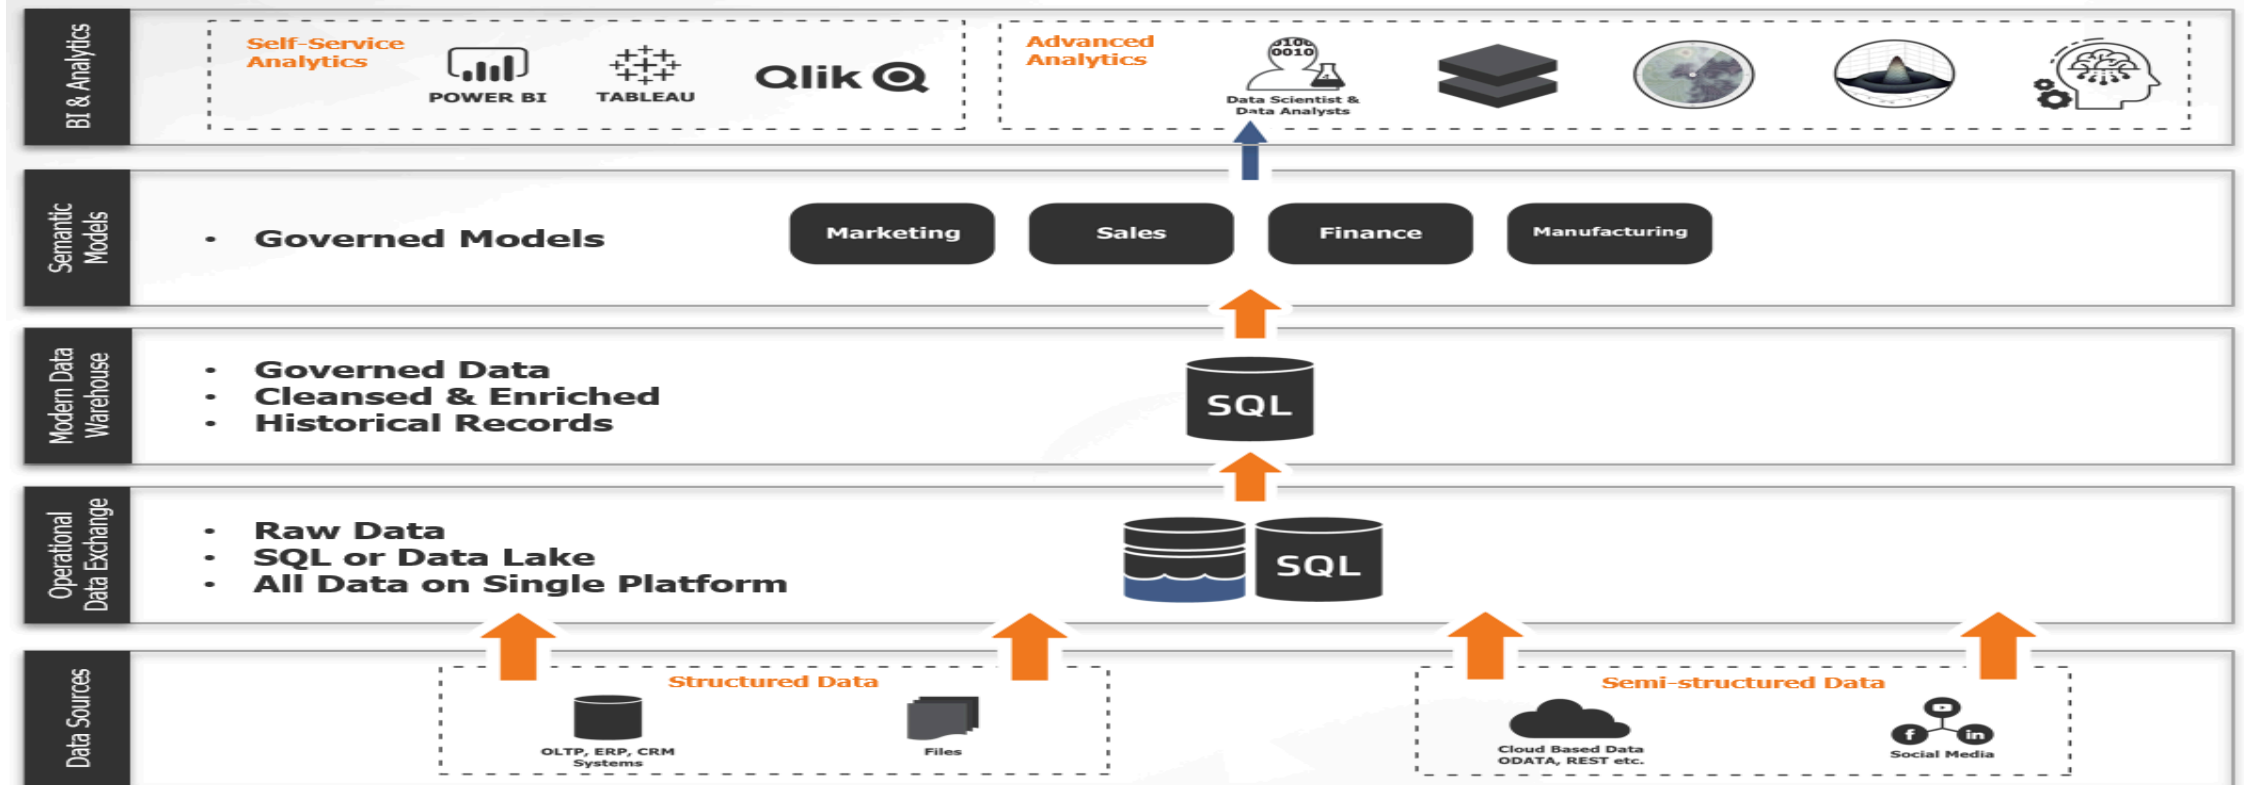
Creator: Upload Audio

-  Submit file + metadata
-  Stored in catalog after validation

Initial Database Architecture

Goal: Scalable, consistent, real-time data processing using only open-source tools.

Approach: Microservice-oriented, data-layered architecture.

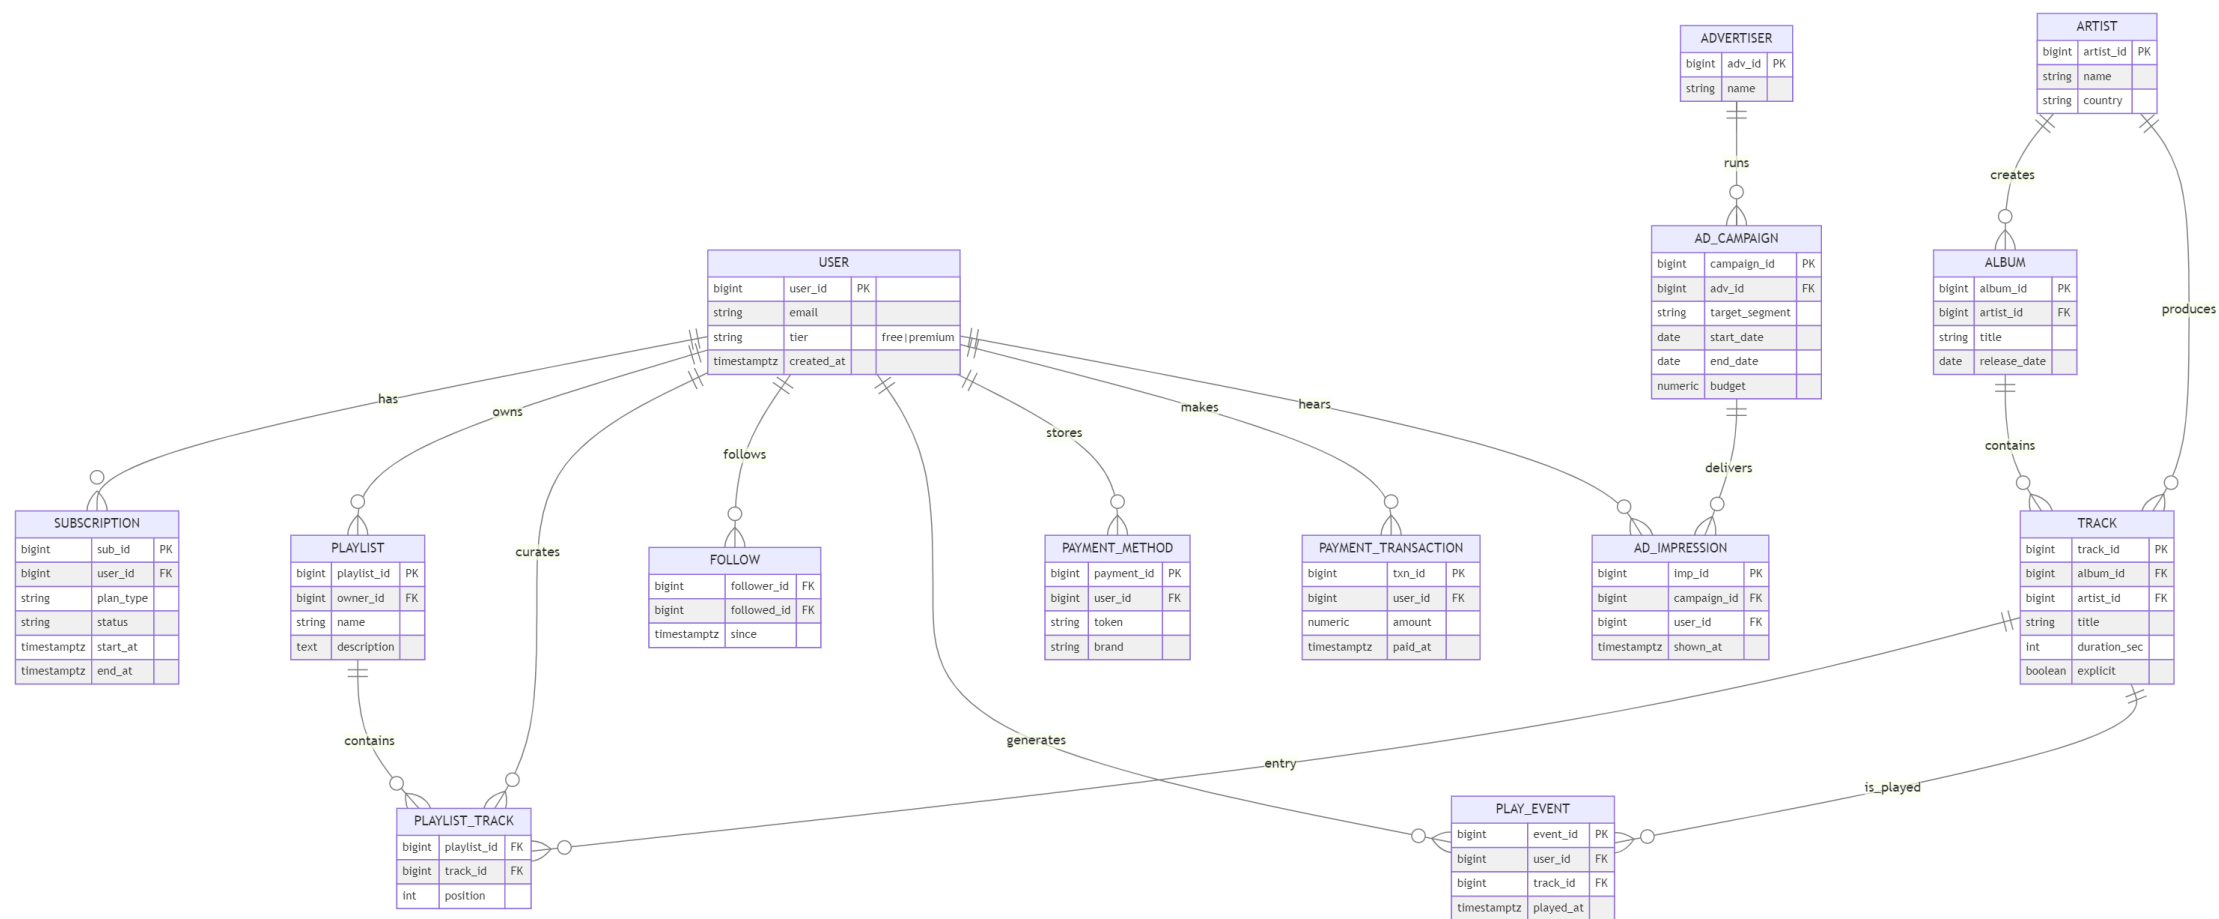




Data Architecture Breakdown

Layer	Purpose	Technologies
Hot OLTP	Critical ops (e.g. billing)	CockroachDB, Citus
Session Store	Playback, telemetry	Cassandra, Scylla
Object Storage	Audio, images	MinIO, Ceph
Search	Full-text autocomplete	OpenSearch
Graph Store	Social graph	Neo4j, JanusGraph
Event Stream	Activity logs	Kafka
Analytics	Real-time + Deep	Flink, Iceberg, Trino

ER Diagram Overview



Conclusion

This project demonstrates:

- ✓ Scalable and distributed data systems
- ✓ Real-time features with Kafka + Flink
- ✓ Secure, compliant design
- ✓ Domain-specific data modeling



References

- Music Business Research (2024)
- Investopedia – *How Spotify Makes Money*
- LinkedIn: *Spotify Tech Stack*
- Intuji – *Spotify Architecture Overview*

 **Thank You!**

Questions?