

# Workshop 2

Data Architecture for a Music Streaming Platform based on Spotify

Brayan Stiven Yate Prada – Student ID: 20192020151

Holman Andres Alvarado Diaz – Student ID: 20201020032

Universidad Distrital Francisco José de Caldas

Faculty of Engineering

Course: Databases II

Professor: Carlos Andrés Sierra Virgüez

May 29, 2025

## Contents

<b>1</b>	<b>Data System Architecture</b>	<b>3</b>
1.1	Justification for Each Technology . . . . .	3
1.2	End-to-End Data Flow (grounded in the schema) . . . . .	4
1.3	Why This Trimmed Stack Meets Our Goals . . . . .	5
<b>2</b>	<b>Information Requirements</b>	<b>5</b>
<b>3</b>	<b>Query Proposals</b>	<b>6</b>
3.1	User and Subscription Management (PostgreSQL) . . . . .	6
3.2	Search and Discovery (OpenSearch) . . . . .	7
3.3	Playback Behaviour (ClickHouse) . . . . .	8
3.4	Campaign and Cohort Analytics (ClickHouse) . . . . .	8
3.5	Royalties and Finance (PostgreSQL) . . . . .	9
<b>4</b>	<b>Improvements Over Workshop 1</b>	<b>10</b>
4.1	Initial Database Architecture . . . . .	10
4.2	Operating Model and Systemic Context . . . . .	10
4.3	Information Flow: Inputs, Outputs and Interactions . . . . .	10
4.4	Critical Bottlenecks Identified . . . . .	11
4.5	System Stress Points . . . . .	11
4.6	Diagnosis of the Central Problem . . . . .	12
4.7	Possible Technical Solutions . . . . .	12
4.8	ER Diagram . . . . .	12
4.9	Overview of the Relational Model . . . . .	12

4.10 Description of Entities and Relationships . . . . .	13
4.11 Storage Strategies and Data Flow . . . . .	14
<b>5 References</b>	<b>14</b>

# 1 Data System Architecture

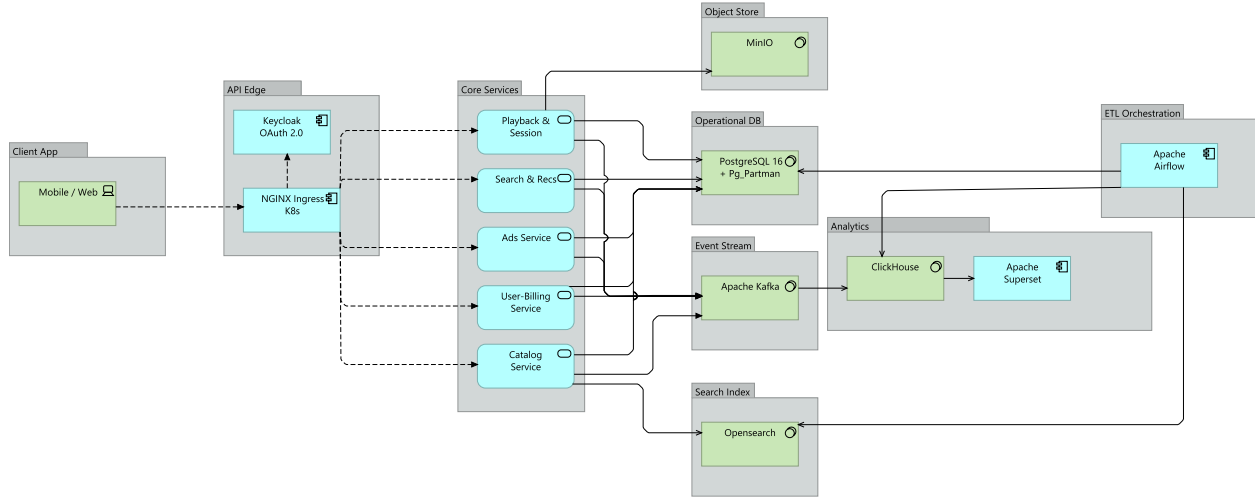


Figure 1: High-level Architecture Diagram

## 1.1 Justification for Each Technology

Component	Technology	Reason	Primary Data Stored
API Edge	<b>NGINX Ingress Controller</b>	Ships with every Kubernetes distribution; trivial TLS termination; rich ecosystem.	Proxies all HTTPS traffic.
Authentication	<b>Keycloak</b>	OAuth2/OIDC server; social log-ins, multi-factor authentication.	User credentials and tokens (hashed / JWT).
Operational SQL	<b>PostgreSQL 16 + pg_partman</b>	Logical/physical replication; partitioning and native JSONB; easy to query.	Users, subscriptions, catalog (artists, albums, tracks), playlists, payments, ads, social graph.
Search / Auto-complete	<b>OpenSearch</b>	Apache-licensed Elasticsearch fork; scales from single node to multi-node; SQL plug-in for analysts.	Denormalised documents for tracks, artists, playlists.
Object Storage	<b>MinIO</b>	Lightweight S3 clone; single binary; erasure coding; deployable on-prem or on any cloud.	Audio files, cover art, long-term raw logs, GDPR exports.

Event Streaming	<b>Apache Kafka</b>	Industry-standard durable, ordered, replayable event log; hundreds of client SDKs.	Play/skip events, ad impressions, change-data capture (CDC) from PostgreSQL via Debezium.
Real-Time / Batch Analytics	<b>ClickHouse</b>	Blazing fast, column-oriented, SQL-native; licence-free; handles > 100 000 inserts/s and sub-second dashboards.	Aggregated play counts, creator dashboards, advertising metrics, Wrapped-style reports.
Orchestration / ETL	<b>Apache Air-flow</b>	Familiar DAG UI; Python-first; vast provider library (Debezium → ClickHouse, backups, ML).	Schedules nightly royalty jobs, table vacuums, data-quality checks.
Business Intelli- gence	<b>Apache Superset</b>	Modern OSS BI; drop-in ClickHouse driver; RBAC and SSO via Keycloak.	Front-end only – reads ClickHouse for charts.

---

## 1.2 End-to-End Data Flow (grounded in the schema)

### Creator Upload

1. The *Catalog Service* receives the audio file and its metadata.
2. Binary data are stored in a MinIO bucket `/audio/`.
3. Rows for artist, album and track are inserted into PostgreSQL (ER diagram tables).
4. A Debezium connector streams these row-level changes to Kafka.
5. A lightweight Flink job flattens the JSON and pushes it to OpenSearch, allowing the track to become searchable within seconds.

### Listener Playback

1. The *Playback Service* retrieves `track_id`, rights and bitrate information from PostgreSQL, signs a presigned MinIO URL and returns it to the client.
2. The client begins streaming; every 5 s it emits a `play-tick` JSON message to Kafka.

**Real-Time Aggregation** ClickHouse materialised views ingest the `play-tick` topic directly—no Spark or Flink cluster is required for the minimum viable product (MVP). Dashboards in Superset refresh automatically for artists and advertising operations.

### Nightly Jobs (*Airflow*)

- `royalty_etl.py`: SQL executed in ClickHouse summarises plays by `track_id` and writes the payout CSV to the PostgreSQL `payouts` table.
- `search_reindex.py`: Performs a bulk re-sync of OpenSearch from PostgreSQL to catch edge cases.
- **Back-ups**: `pg_dump` streams are stored as compressed objects in MinIO.

**Social Graph and Recommendations** The `user_follow` bridge table lives in PostgreSQL; for the MVP we run simple collaborative-filtering in Python inside Airflow once per day and write top- $N$  recommendations into `user_recs`. When sub-50 ms cold-start latencies become necessary these features can migrate to Redis, but not in version 1.

### 1.3 Why This Trimmed Stack Meets Our Goals

- **Feasible for a small team:** only six stateful services (PostgreSQL, MinIO, Kafka, OpenSearch, ClickHouse, Keycloak). All ship Helm charts and run on three modest VMs or a single Kubernetes cluster.
- **Single source of truth:** PostgreSQL hosts the entire ER diagram with just two bridge tables (`playlist_track`, `user_follow`) and one fact table (`ad_impression`). OLTP rows are partitioned by tenant (`user_id mod N`) enabling future shard-out.
- **Analytics without Hadoop:** ClickHouse reads Kafka directly, avoiding Spark/HDFS.
- **Search within seconds:** Debezium + Kafka keeps OpenSearch in near-real-time sync.
- **Fully open-source:** no licences, no managed-only products.
- **Clear upgrade path:** if traffic explodes we can add Postgres replicas, expand Kafka, deploy more ClickHouse shards or adopt Citus; none require a rewrite.

## 2 Information Requirements

#	Information Type	Description	Storage	Business Link	Key User Stories
1	<b>Catalog Metadata</b>	Track, album, artist info. Enables search & legal playback.	PostgreSQL ( <code>tracks</code> , <code>albums</code> , <code>artists</code> ); OpenSearch cache.	Global audio library; Key Partners.	Search; Playback; Upload Audio.
2	<b>Audio References</b>	Presigned URLs + bitrate list for playback via CDN or MinIO.	Generated by Playback Service from object metadata.	Seamless streaming.	Playback; Device Integration.
3	<b>User Profile</b>	Display name, avatar, language, tier, parental controls. UI personalization + policy enforcement.	PostgreSQL ( <code>users</code> ).	Personalisation; Segmentation.	Profile Management; Register; Parental Control.
4	<b>Subscription &amp; Billing</b>	Plan, renewal, token, grace state. Controls privileges and ad access.	PostgreSQL ( <code>subscriptions</code> , <code>payment_method</code> ).	Recurring revenue.	Subscription Management; Premium Access.

5	<b>Playback State</b>	Position, device ID, network, heartbeat. Enables adaptive bitrate + resume.	Kafka → ClickHouse or in-memory cache.	Playback continuity.	Adaptive Streaming; Resume Playback.
6	<b>Suggestion</b>	Ranked track IDs + labels (e.g., repeat, trending).	PostgreSQL JSONB ( <code>user_recs</code> ).	Discovery; ML-driven value.	Discover; Playlist Creation.
7	<b>Search Suggestions</b>	Tokens, fuzzy matches (tracks, artists, playlists).	OpenSearch.	Search efficiency.	Content Search.
8	<b>Social Graph</b>	Follows, playlist rights, friend activity.	PostgreSQL ( <code>user_follow</code> , <code>playlist_track</code> ); Kafka.	Community features.	Follow Users; Collaborative Playlists.
9	<b>Ads &amp; Creatives</b>	Targeting, CPM, assets; user context.	PostgreSQL ( <code>ad_campaign</code> , <code>ad_creative</code> ) + in-memory rules.	Ad revenue.	Launch Campaigns; Real-Time Ads.
10	<b>Creator Analytics</b>	Play counts, geo maps, promo-lift.	ClickHouse views via Superset.	Creator tools; Partner support.	View Stats; Track Promotion Impact.
11	<b>Royalty Reports</b>	Stream counts, payouts per rightsholder.	Airflow ETL → PostgreSQL ( <code>payouts</code> ); MinIO.	Royalty costs.	(Internal workflows).
12	<b>Compliance Logs</b>	Admin actions, GDPR links, token audits.	Append-only Postgres + WORM MinIO.	Trust; Legal compliance.	Data Requests; Audit Trails.

---

## Traceability to Business Model and User Stories

- **Value Proposition** (“global audio library”, ad-free Hi-Fi, discovery) requires fast access to items 1, 2, 6 and 5.
- **Revenue Streams** (subscriptions & ads) rely on items 4 and 9.
- **Customer Relationships / Segments** (community, stats) depend on items 3, 8 and 10.
- **Cost Structure / Key Partners** (royalty payouts) require items 1 and 11.
- **Compliance & Trust** are covered by item 12 and secured payment tokens.

## 3 Query Proposals

### 3.1 User and Subscription Management (PostgreSQL)

-- 1A. Premium users in top-5 high-usage countries

```

WITH top_countries AS (
    SELECT country
    FROM   play_events_daily -- daily roll-up loaded by Airflow
    ORDER BY total_plays DESC
    LIMIT  5
)
SELECT u.user_id,
       u.display_name,
       s.plan,
       s.country
FROM   users u
JOIN   subscriptions s USING (user_id)
WHERE  s.status      = 'active'
      AND s.plan      = 'premium'
      AND s.country  IN (SELECT country FROM top_countries);

```

Purpose: identify premium markets for targeted marketing.

```

-- 2A. Potential account sharing (>=3 countries in 30 days)
SELECT user_id,
       COUNT(DISTINCT ip_country) AS unique_countries_30d
FROM   session_logs PARTITION FOR (CURRENT_DATE - INTERVAL '30 days')
GROUP BY user_id
HAVING COUNT(DISTINCT ip_country) > 3;

```

Purpose: detect family-plan abuse.

## 3.2 Search and Discovery (OpenSearch)

```

GET /tracks/_search
{
  "query": {
    "function_score": {
      "query": { "match": { "title": "drake" } },
      "field_value_factor": {
        "field": "play_count",
        "factor": 0.1,
        "modifier": "sqrt",
        "missing": 1
      },
      "boost_mode": "multiply"
    }
  },
  "sort": [
    { "_score": "desc" },
    { "release_date": "desc" }
  ]
}

```

```
]
}
```

Purpose: relevancy-boosted autocomplete.

### 3.3 Playback Behaviour (ClickHouse)

-- 3A. Early-churn indicator for new sign-ups (<=30s skips)

```
SELECT
    user_id,
    COUNT()          AS total_plays,
    sum(duration < 30) AS short_skips
FROM   play_events
WHERE  signup_date >= today() - 7
GROUP BY user_id
HAVING total_plays > 5;
```

-- 3B. Rapid-skip bot detection (last 10 min)

```
SELECT *
FROM   play_events
PREWHERE user_id = 'u_999'
        AND event_time > now() - INTERVAL 10 MINUTE
WHERE  duration < 15;
```

### 3.4 Campaign and Cohort Analytics (ClickHouse)

-- 4A. Retention by sign-up cohort

```
WITH cohort AS (
    SELECT user_id,
           toDate(min(event_time)) AS cohort_date
    FROM   play_events
    GROUP BY user_id
)
SELECT cohort_date,
       activity_date,
       uniqExact(user_id) AS active_users
FROM (
    SELECT user_id,
           toDate(event_time) AS activity_date
    FROM   play_events
) e
JOIN cohort USING user_id
GROUP BY cohort_date, activity_date
ORDER BY cohort_date, activity_date;
```



```

-- 4B. Streams attributable to promotional campaigns
SELECT
    t.artist_id,
    t.track_id,
    c.campaign_id,
    count() AS plays_during_campaign
FROM   play_events      AS pe
JOIN   tracks            AS t  ON pe.track_id = t.track_id
JOIN   ad_campaign       AS c  ON pe.event_time BETWEEN c.start_date AND c.end_date
WHERE  c.campaign_type = 'promotional'
GROUP BY
    t.artist_id, t.track_id, c.campaign_id;

```

### 3.5 Royalties and Finance (PostgreSQL)

```

-- 5A. Royalty calculation weighted by local rates
SELECT
    t.artist_id,
    t.track_id,
    pe.country,
    COUNT(*) * r.rate AS royalty_amount
FROM   play_events_monthly pe -- ETL roll-up table
JOIN   tracks            t  ON pe.track_id = t.track_id
JOIN   royalty_rates      r  ON r.country = pe.country
                                AND r.artist_id = t.artist_id
WHERE  pe.month = '2025-05'
GROUP BY t.artist_id, t.track_id, pe.country, r.rate;

-- 5B. Cross-check: plays vs payments divergence >1 %
WITH plays AS (
    SELECT artist_id,
           COUNT(*) AS total_plays
    FROM   play_events_monthly
    WHERE  month = '2025-05'
    GROUP BY artist_id
)
SELECT p.artist_id,
       SUM(ap.royalty_amount) AS total_paid,
       total_plays
FROM   artist_payouts ap
JOIN   plays            p USING (artist_id)
GROUP BY p.artist_id, total_plays
HAVING ABS(SUM(ap.royalty_amount) / NULLIF(total_plays,0) - expected_rate) > 0.01;

```

## Technology Mapping

Use-Case	Query Engine	Rationale
Operational joins, financial exactness	PostgreSQL	ACID semantics and referential integrity.
Event-stream ad-hoc, cohort, fraud analysis	ClickHouse	100 000 RPS ingestion and sub-second scans.
Search / autocomplete	OpenSearch	Token scoring and fuzzy match.
Event ingestion	Kafka	Decouples writes and enables real-time pipelines.

## 4 Improvements Over Workshop 1

### 4.1 Initial Database Architecture

Designing a database architecture for a large-scale music-streaming platform requires a balance of technical and business constraints: ultra-low playback latency, strict compliance with regulations such as GDPR, global scalability and financial integrity. The architecture therefore distributes data into specialised layers, selecting *only* open-source technologies to ensure flexibility, transparency and long-term sustainability.

### 4.2 Operating Model and Systemic Context

Spotify adopts a *freemium* model: free users are monetised through advertising while Premium users pay recurring subscriptions. Roughly 70 % of revenue is allocated to royalty payments for artists and record labels.

The data ecosystem involves multiple actors:

- End users (mobile / web clients)
- Backend and distributed storage layer
- Content creators
- Advertising subsystem
- Analytics and reporting platform
- Contractual royalty rules

### 4.3 Information Flow: Inputs, Outputs and Interactions

The system operates via asynchronous, bidirectional data flows. Principal routes are:

#### Inputs

Multimedia files (audio, cover art, metadata); user events (play, like, follow); financial information and credentials.

#### Outputs

Optimised music playback; personalised recommendations; analytical reports; real-time targeted advertisements.

## Interactions

Client → Application → Microservice backend → Distributed storage, with lateral interactions to Advertising and Recommendation subsystems.

## 4.4 Critical Bottlenecks Identified

Subsystem	Problem	Risk	Recommendations
Intensive social-graph reads	Notable latencies (P99 > 200 ms)	User experience degradation	Partition graph, introduce specialised cache (e.g. Dgraph or Redis-Graph).
Advertising	Cache saturation under high concurrency	Ad-decision latency > 50 ms	Hierarchical LRU cache, pre-computed segments.
Streaming	Real-time CDN and DRM selection	Start time > 300 ms causes abandonment	Edge nodes with pre-signed URLs; regional CDN hints.
Royalties	Massive event calculation	Legal/reputational risk	Incremental aggregation, ClickHouse materialised views, parallel pipelines.
GDPR	Data proliferation	Non-compliance fines	Unified subject-ID mapping and automated data-subject export routines.

## 4.5 System Stress Points

Potential Crisis	Technical and Business Cause
Explosive concurrent-user growth	Recommendation, playback or scaling failures due to horizontal stress.
Advertising-decision delays	Revenue loss and poor user experience for free tier.
Inaccurate royalty calculations	Legal and reputational consequences with rights holders.
Content-ingestion saturation (10 TB / day)	Bottlenecks in validation, ML pipelines and media persistence.
Social-graph synchronisation errors	Inconsistent likes/follows (latency > 5 s).
Autoscaling failures	Frozen experience due to unbalanced or mis-distributed services.

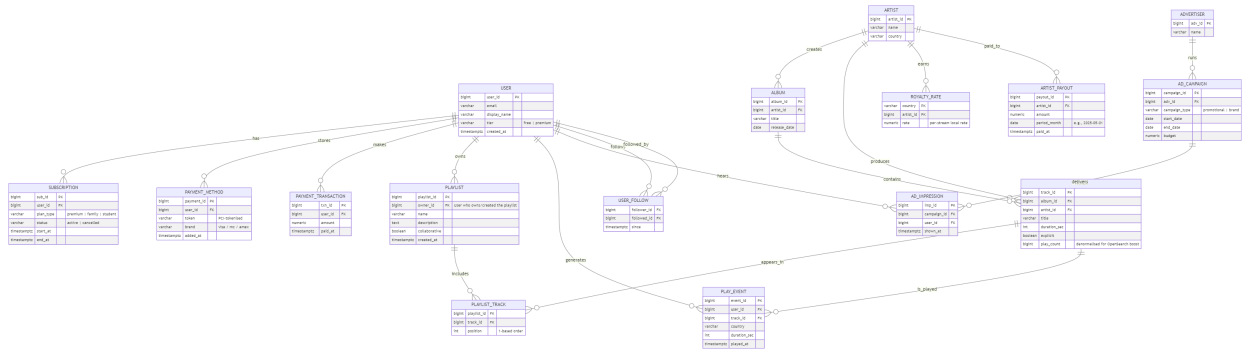


Figure 2: Caption

## 4.6 Diagnosis of the Central Problem

Spotify faces constant tension among horizontal scalability, ultra-low latency and contractual compliance: every new user, artist or advertising campaign represents a potential distributed load hotspot.

## 4.7 Possible Technical Solutions

1. Strengthen distributed cache for personalization and advertising.
2. Apply edge-computing for initial playback near the client.
3. Decouple royalty pipelines into asynchronous, parallel layers.
4. Deploy full distributed tracing (OpenTelemetry).
5. Optimize graph-partitioning algorithms and advertising clusters.

**Update:** A systemic problem analysis has been added, identifying key points, bottlenecks, information flows and a clearer operating model, clarifying the issue to solve and updating the ER diagram accordingly.

## 4.8 ER Diagram

## 4.9 Overview of the Relational Model

The proposed relational model structures the critical information of a musical-content, advertising and social-interaction platform, based on a microservices architecture designed for large-scale operation. The implementation relies on PostgreSQL 16, taking advantage of referential integrity, ACID transactions, advanced types (e.g. JSONB) and native sharding (via `pg_partman` or Citus).

The model is organised into six functional domains, each with independent entities and clear relationships, enabling horizontal scalability, auditability and extensibility in line with distributed-design principles and regulatory compliance.

## 4.10 Description of Entities and Relationships

### User and Subscription Domain

- **USER**: end-user profiles, including unique identifier, subscription type (free / premium) and creation date.
- **SUBSCRIPTION**: transactional table storing the history of active and cancelled plans per user.
- **PAYMENT\_METHOD** and **PAYMENT\_TRANSACTION**: manage tokenised payment methods (PCI-DSS) and related transactions, linked to **USER** via foreign keys.

This flow provides complete traceability of the account life-cycle and its monetisation.

### Catalog Domain

- **ARTIST**, **ALBUM** and **TRACK**: compose the musical catalogue hierarchy. **TRACK** includes key metadata such as duration, explicit-content flag and a denormalised `play_count` field to improve analytical indexing in OpenSearch.
- Relationships are modelled *one-to-many* from artist to album and album to track.

### Playlists and Social Domain

- **PLAYLIST**: associated with a user (`owner_id`) and supporting collaborative lists.
- **PLAYLIST\_TRACK**: many-to-many bridge table preserving sequential order by `position` column.
- **USER\_FOLLOW**: represents the social follow graph, using composite keys and timestamps for temporal tracking.

These entities enable social functionality and personalisation with low logical coupling.

### Playback and Events Domain

- **PLAY\_EVENT**: records each playback session with contextual metadata such as country, effective duration and timestamp. Modelled as a very high-cardinality event table, suitable for replication into OLAP or time-series systems.

### Advertising Domain

- **ADVERTISER**, **AD\_CAMPAIGN**, **AD\_IMPRESSION**: model advertising campaigns and their interaction with users.
- **AD\_IMPRESSION** functions as a fact table, registering which user was exposed to which campaign and when.

### Royalties and Finance Domain

- **ROYALTY\_RATE**: defines the per-play fee by artist and country.
- **ARTIST\_PAYOUT**: represents consolidated monthly payments to each artist.

Both entities are populated from aggregated **PLAY\_EVENT** data, enabling contractual and regional revenue calculations with full financial traceability.

## 4.11 Storage Strategies and Data Flow

- Implemented on PostgreSQL 16 with horizontal sharding by `user_id` in high-volume tables (`PLAY_EVENT`, `USER_FOLLOW`, `AD_IMPRESSION`).
- Events are written to PostgreSQL and synchronised to external analytical stores (ClickHouse, OpenSearch) for real-time exploitation.
- Sensitive tables such as `PAYMENT_METHOD` integrate external tokenisation and at-rest encryption (AES-256).
- Retention, auditing (`created_at`, `paid_at`, `shown_at`) and strong consistency policies are enforced on critical tables (`SUBSCRIPTION`, `TRACK`, `ARTIST_PAYOUT`).

## 5 References

### References

- [1] Corporate Finance Institute. *Business Model Canvas Examples*. Available at: <https://corporatefinanceinstitute.com/resources/management/business-model-canvas-examples/>
- [2] Business Model Analyst. *Spotify Business Model*. Available at: <https://businessmodelanalyst.com/spotify-business-model/>
- [3] Music Business Research. (19 August 2024). *The Music Streaming Economy Part 10: Spotify's Business Model*. Available at: <https://musicbusinessresearch.wordpress.com/2024/08/19/the-music-streaming-economy-part-10-spotifys-business-model/>
- [4] Investopedia. *How Spotify Makes Money*. Available at: <https://www.investopedia.com/articles/investing/120314/spotify-makes-internet-music-make-money.asp>
- [5] IIDE. *Business Model of Spotify*. Available at: <https://iide.co/case-studies/business-model-of-spotify/>
- [6] GrowthX Club. *Spotify Business Model*. Available at: <https://growthx.club/blog/spotify-business-model#spotify-revenue-model>
- [7] Software Developer Diaries. *How Spotify's Playback Works Under the Hood*. YouTube video. Available at: <https://www.youtube.com/watch?v=K26bGXVR-mE>
- [8] Lopez, M. A. *Deep Dive into the Tech Stack Used by Spotify*. LinkedIn article. Available at: <https://www.linkedin.com/pulse/deep-dive-tech-stack-used-spotify-marny-a-lopez>
- [9] Intuji. *How Does Spotify Work? – Tech Stack Explored*. Available at: <https://intuji.com/how-does-spotify-work-tech-stack-explored/>